



Introduction to Apache Spark

Fernando Cores

Index

- + 3.1. What is Spark?
- + 3.2. The Spark Programming model
- + 3.3. Working with Resilient Distributed Datasets (RDDs)
- + 3.4. Programming with Spark

How to access to Spark (I)

- + We are going to use the Spark installation in the Pirgi big-data cluster (pirgi.udl.cat)
- + You can work with Spark in two different ways:
 - Using Jupyter Notebooks through a web-browser in the following page:
 - <http://pirgi.udl.cat:8000>
 - You have to login with your pirgi's account.
 - Using pyspark shell or launching yarn spark jobs:
 - You have to connect to the pirgi cluster using ssh:
 - `ssh username@pirgi.udl.cat`
 - and execute the pyspark shell:
 - `pyspark --master yarn-client --driver-memory 1g --executor-memory 1g --executor-cores 1`
 - or launch the spark yarn job:
 - `spark-submit --master yarn --deploy-mode cluster --driver-memory 1g --executor-memory 512m --executor-cores 1 ./WordCount.py`

How to access to Spark (Docker I)

- + You can also use a Jupyter Spark image: jupyter/all-spark-notebook (5.26 GB)
 - Spark 2.0.2 with Hadoop 2.7, Jupyter Notebook 4.3.x, Python 3.x and Python 2.7.x environments
 - Available: <https://github.com/jupyter/docker-stacks/tree/master/all-spark-notebook>
- + Pull the image:
 - `docker pull jupyter/all-spark-notebook latest`
- + Running the docker machine:
 - `docker run -it --rm -p 8888:8888 -v<local_directory>:/home/jovyan/work -w /home/jovyan/work jupyter/all-spark-notebook`
- + Login on Web browser with the generated token:
`http://localhost:8888/?token=931d86475of56f8ec301e9ed4525a4665f8869ageb718ecc`

How to access to Spark (Docker II)

- + You can also use a docker standalone Apache Spark (2.88 GB)
 - Hadoop 2.6.0 and Apache Spark v1.6.0 on Centos
 - Available: <https://hub.docker.com/r/sequenceiq/spark/>
- + Pull the image:
 - `docker pull sequenceiq/spark:1.6.0`
- + Running the docker machine:
 - `docker run -it -p 8088:8088 -p 8042:8042 -h sandbox sequenceiq/spark:1.6.0 bash`



What is Spark?

What is Spark?




- + An Apache Foundation **open source** project; not a product
- + Enables highly **iterative** analysis on **massive** volumes of data at scale
- + An **in-memory computing** engine that works with **distributed** data; not a data store
- + **Unified environment** for data scientists, developers and data engineers
- + Radically simplifies the process of developing **intelligent apps** fuelled by data

Spark Motivation

Current popular programming models for clusters transform data flowing from stable storage to stable storage.

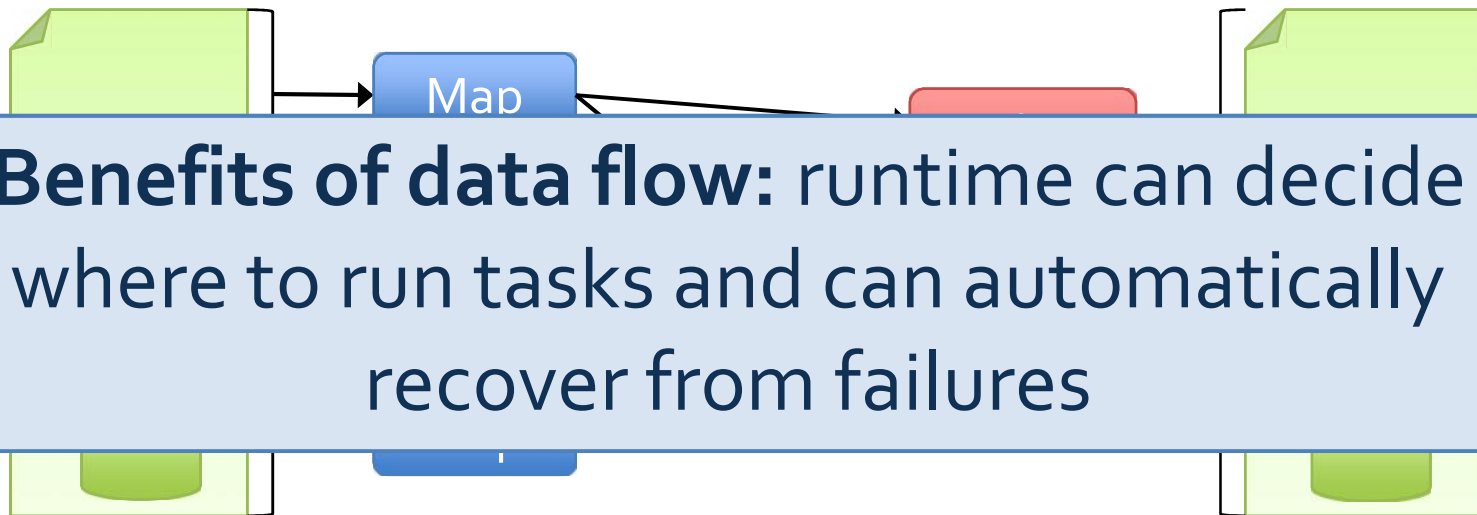
Example: MapReduce:



Spark Motivation

Current popular programming models for clusters transform data flowing from stable storage to stable storage.

Example: MapReduce:



Benefits of data flow: runtime can decide where to run tasks and can automatically recover from failures

Spark Motivation

- + Acyclic data flow is a powerful abstraction, but is not efficient for applications that repeatedly reuse a *working set* of data:
 - **Iterative** algorithms (many in machine learning)
 - **Interactive** data mining tools (R, Excel, Python)
- + Spark makes working sets a first-class concept to efficiently support these apps

Spark Goal

- + Provide distributed memory abstractions for clusters to support apps with working sets
- + Retain the attractive properties of MapReduce:
 - Fault tolerance (for crashes & stragglers)
 - Data locality
 - Scalability

Solution: augment data flow model with “resilient distributed datasets” (RDDs)

Apache Spark is...

Fast


- Leverages aggressively cached in-memory distributed computing and JVM threads
- Faster than MapReduce for some workloads

Ease of use (for programmers)


- Written in Scala, an object-oriented, functional programming language
- Scala, Python and Java APIs
- Scala and Python interactive shells
- Runs on Hadoop, Mesos, standalone or cloud

General purpose

- Covers a wide range of workloads
- Provides SQL, streaming and complex analytics

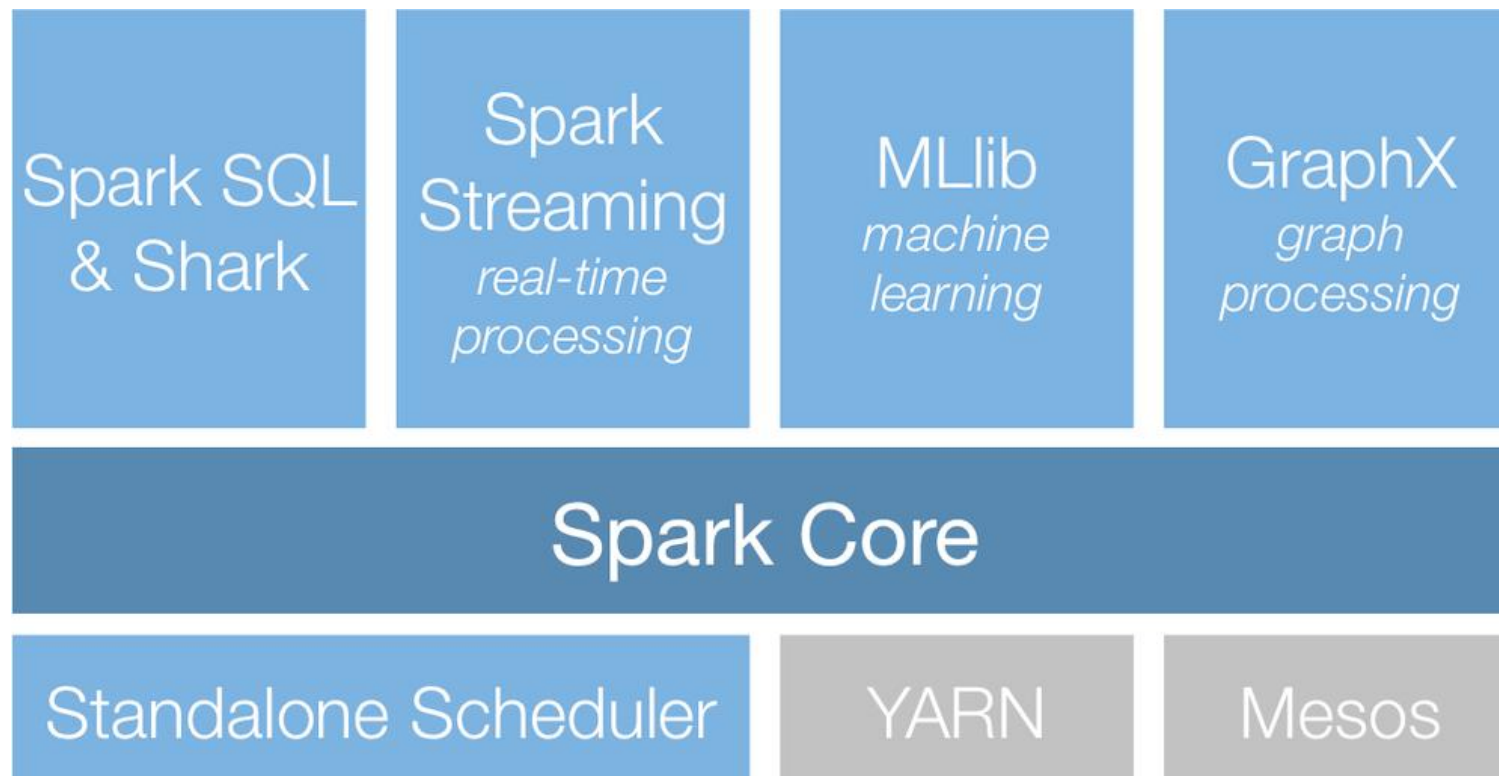


Logistic regression in Hadoop and Spark



from <http://spark.apache.org>

Spark Stack




The Spark Programming model



Programming Model


- + Resilient distributed datasets (RDDs)
 - Immutable collections partitioned across cluster that can be rebuilt if a partition is lost
 - Created by transforming data in stable storage using data flow operators (map, filter, group-by, ...)
 - Can be *cached* across parallel operations
- + Parallel operations on RDDs
 - Reduce, collect, count, save, ...
- + Restricted shared variables
 - Accumulators, broadcast variables

Spark Application (I)



- + A Spark application consists of a driver program that launches various parallel operations on a cluster.
 - The driver program contains application's main function and defines distributed datasets on the cluster, then applies operations to them.


Spark Application (II)



- + Driver programs access Spark through a **SparkContext** object.
 - Represents a connection to a computing cluster. It is used to create and manipulate distributed datasets and shared variables
 - **SparkContext** is initialized with an instance of a **SparkConf** object, which contains various Spark cluster-configuration settings.

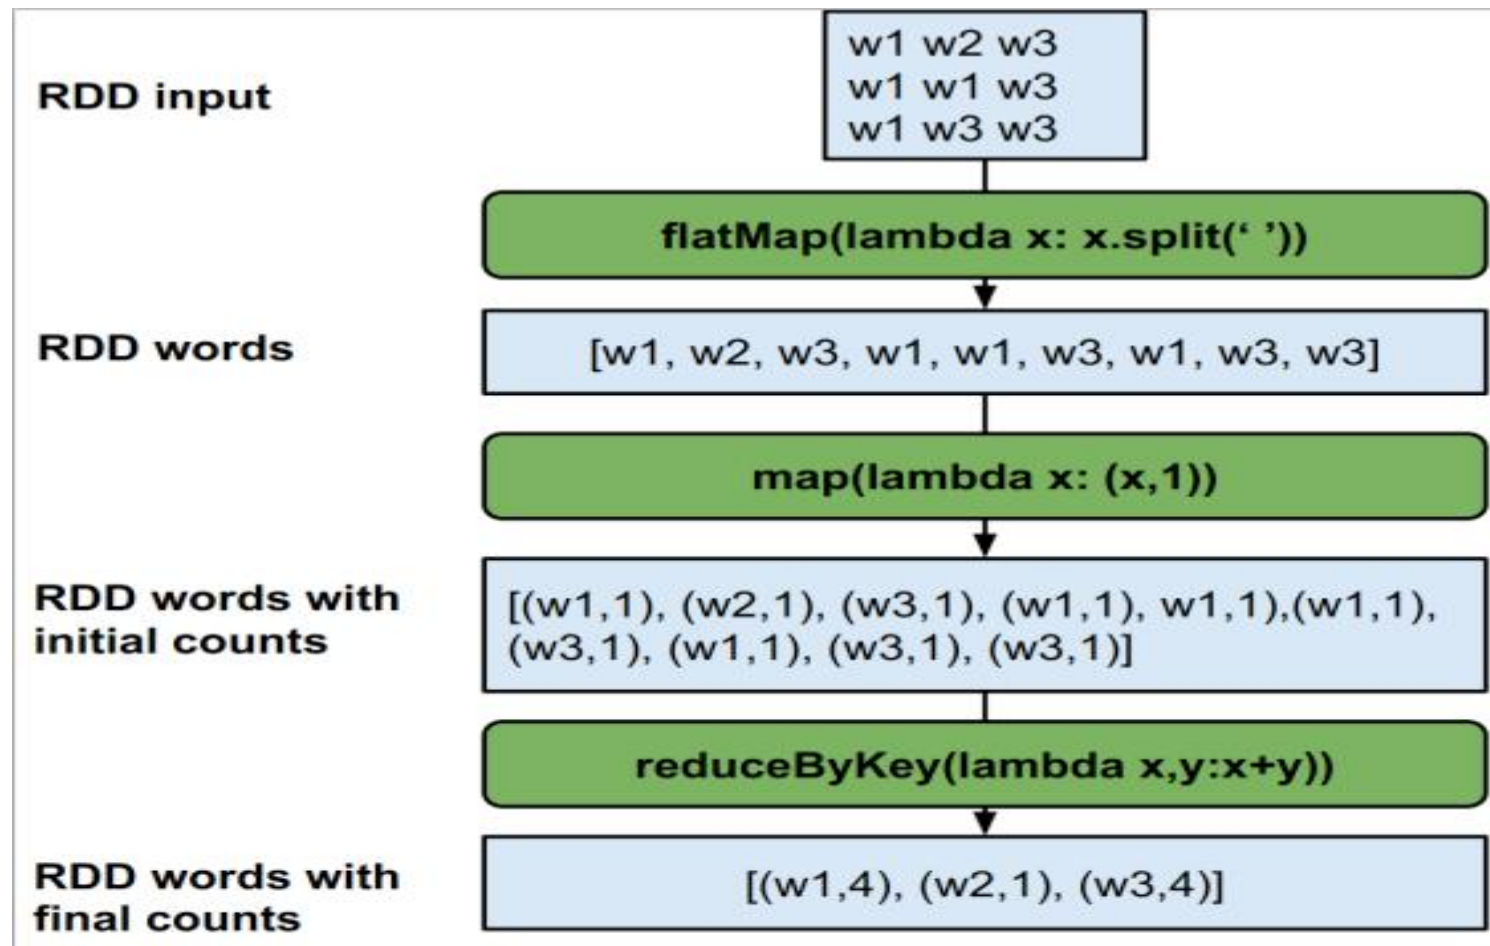
```
from pyspark import SparkConf, SparkContext
conf = SparkConf().setMaster("local").setAppName("My App")
sc = SparkContext(conf = conf)
```

Spark Application (III)



- + To run their operations, driver programs typically manage a number of nodes called **executors**.
 - They perform individual tasks and returns the results to Driver
 - Provide memory storage for datasets

Example: WordCount in Spark (I)



Example: WordCount in Spark (II)

Python

```
from pyspark import SparkContext
sc = SparkContext(appName="WordCount Spark App")
text_file = sc.textFile("hdfs://...")
counts = text_file.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)
counts.saveAsTextFile("hdfs://...")
counts.count()
```

Scala

```
val sc = new SparkContext("local[2]", "WordCount Spark App")
val textFile = sc.textFile("hdfs://...")
val counts = textFile.flatMap(line => line.split(" "))
    .map(word => (word, 1))
    .reduceByKey(_ + _)
counts.saveAsTextFile("hdfs://...")
counts.count()
```

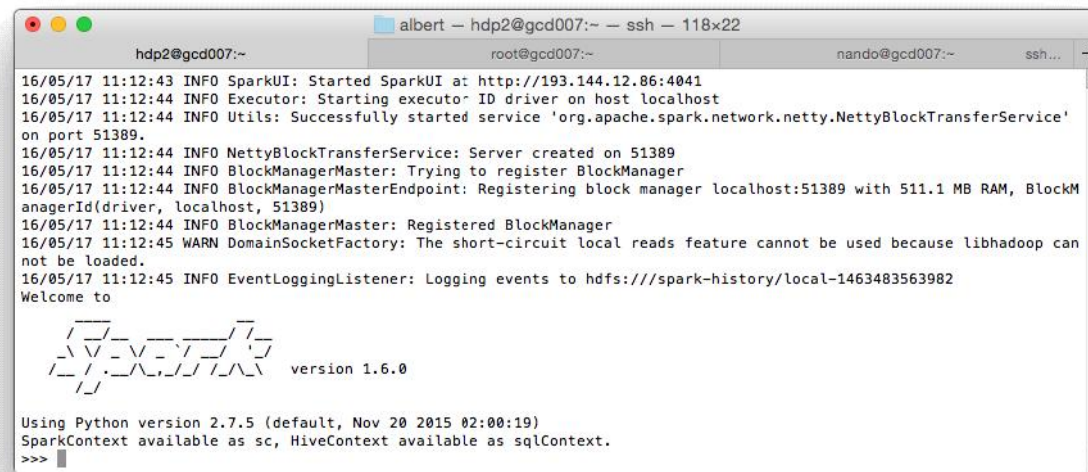
Example: WordCount in Spark (III)

Java

```
public class JavaApp {  
    public static void main(String[] args) {  
        JavaSparkContext sc = new JavaSparkContext("local[2]", "WordCount Spark App");  
        JavaRDD<String> textFile = sc.textFile("hdfs://...");  
        JavaRDD<String> words = textFile.flatMap(new FlatMapFunction<String, String>() {  
            public Iterable<String> call(String s) { return Arrays.asList(s.split(" ")); }  
        });  
        JavaPairRDD<String, Integer> pairs = words.mapToPair(new PairFunction<String, String, Integer>() {  
            public Tuple2<String, Integer> call(String s) { return new Tuple2<String, Integer>(s, 1); }  
        });  
        JavaPairRDD<String, Integer> counts = pairs.reduceByKey(new Function2<Integer, Integer, Integer>() {  
            public Integer call(Integer a, Integer b) { return a + b; }  
        });  
        counts.saveAsTextFile("hdfs://...");  
    }  
}
```

Spark's Shells

- + The Spark shell provides an easy and convenient way to quickly prototype certain operations without having to develop a full program, packaging it and then deploying it.
- + Spark supports writing programs interactively using either the Scala or Python REPL (interactive shell).
 - Scala shell: `./bin/spark-shell`
 - Python shell: `./bin/pyspark`



```
albert - hdp2@gcd007:~ - ssh - 118x22
hdp2@gcd007:~
16/05/17 11:12:43 INFO SparkUI: Started SparkUI at http://193.144.12.86:4041
16/05/17 11:12:44 INFO Executor: Starting executor ID driver on host localhost
16/05/17 11:12:44 INFO Utils: Successfully started service 'org.apache.spark.network.netty.NettyBlockTransferService'
on port 51389.
16/05/17 11:12:44 INFO NettyBlockTransferService: Server created on 51389
16/05/17 11:12:44 INFO BlockManagerMaster: Trying to register BlockManager
16/05/17 11:12:44 INFO BlockManagerMasterEndpoint: Registering block manager localhost:51389 with 511.1 MB RAM, BlockM
anagerId(driver, localhost, 51389)
16/05/17 11:12:44 INFO BlockManagerMaster: Registered BlockManager
16/05/17 11:12:45 WARN DomainSocketFactory: The short-circuit local reads feature cannot be used because libhadoop can
not be loaded.
16/05/17 11:12:45 INFO EventLoggingListener: Logging events to hdfs:///spark-history/local-1463483563982
Welcome to
  ____
 /_  __ \
/_/_/  \_/_/
version 1.6.0

Using Python version 2.7.5 (default, Nov 20 2015 02:00:19)
SparkContext available as sc, HiveContext available as sqlContext.
>>>
```

Practice: Execute WordCount in the pyspark shell

+ Execute the WordCount spark application in the pyspark shell

- Open the shell

➤ `/bin/pyspark`

- Use the books directory as input:

- `hdfs:///shared/nando/data/books/`

- Copy and execute each line of the code

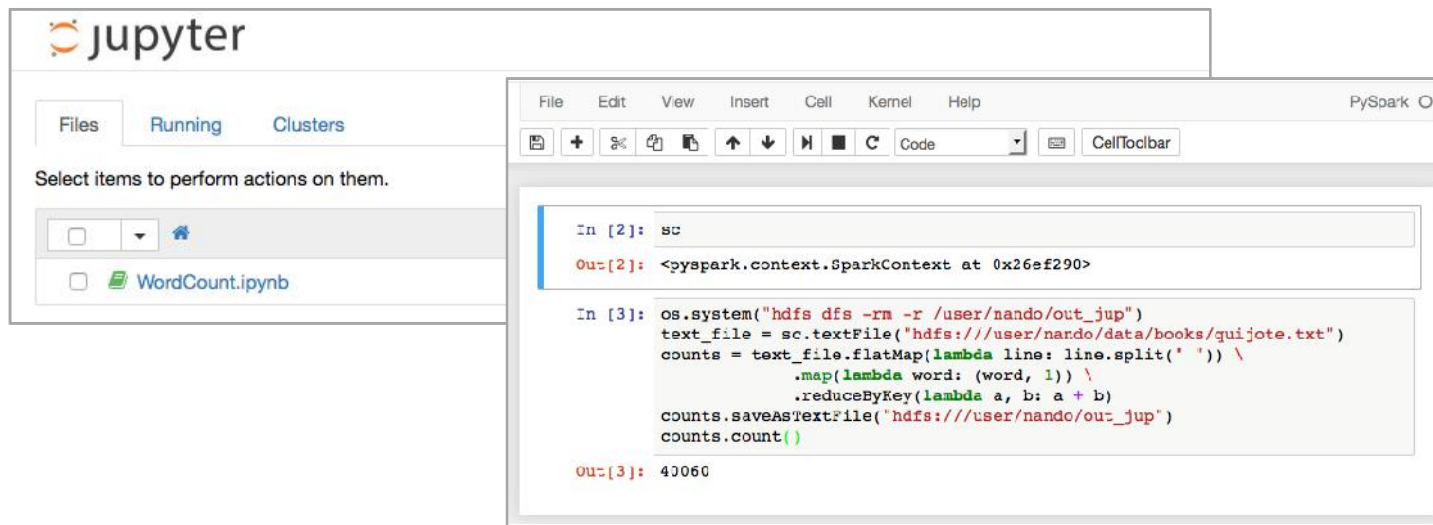
```
os.system("hdfs dfs -rm -r /shared/nando/output/spark/wc")
text_file = sc.textFile("hdfs:///shared/nando/data/books/")
counts = text_file.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)
counts.saveAsTextFile("hdfs:///shared/nando/output/spark/wc")
counts.count()
```

- Check the results:

- `hdfs dfs -cat /user/<your_user>/<out_dir>/p*`

Jupyter Notebook

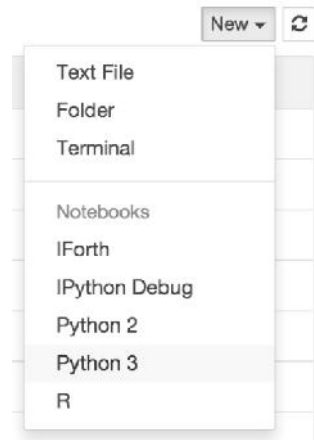
- + **Jupyter** notebook is an interactive Python shell which lets you interact with your data one step at a time and also perform simple visualizations
 - Supports tab auto-completion on class names, functions, methods, variables.
 - Offers more explicit and colour-highlighted error messages than the command line python shell.
 - Provides integration with basic UNIX shell allowing to run simple shell commands.




Jupyter Notebook Basics (I)



- + The notebook dashboard serves as a home page for the notebook. Its main purpose is to display the notebooks and files in the current directory.
 - The top of the notebook list displays clickable contents of the current directory.
 - By clicking on these element or on sub-directories in the notebook list, you can navigate your file system.
 - It also allows the creation of new directories, upload files or run notebooks.
 - To shutdown, delete, duplicate, or rename a notebook check the checkbox next to it and an array of controls will appear at the top of the notebook list
 - To create a new notebook, click on the "New" button at the top of the list and select a kernel from the dropdown.



Jupyter Notebook Basics (II)




- + The notebook user interface (UI) allows you to run code and author notebook documents interactively.
 - The notebook UI has the following main areas: Menu, Toolbar and Notebook area and cells
 - The notebook has a modal user interface. This means that the keyboard does different things depending on which mode the Notebook is in.
 - There are two modes: **edit mode** and **command mode**.
 - There are two main types of cells: Code and Markdown cells
 - Code cells allows to write and execute spark code
 - Markdown cells allows to write markup text, Markdown, that it is a superset of HTML

```
1. Practice: Execute WordCount in the Notebook
* Input Data: Datasets/Books
* Result: Results/Out_WordCount
```

```
1. Practice: Execute WordCount in the Notebook
• Input Data: Datasets/Books
• Result: Results/Out_WordCount
```

```
In [2]: import pyspark
        sc = pyspark.SparkContext('local[*]')
```

Jupyter Notebook Basics (III)



+ The Jupyter Notebook is an interactive environment for writing and running code. The notebook is associated with the IPython kernel, therefore runs Python code.

- Code cells allow you to enter and run code
 - Alt-Enter runs the current cell and inserts a new one below.
 - Ctrl-Enter run the current cell and enters command mode.

```
In [2]: a = 10
```

```
In [3]: print(a)
```

```
10
```

- Code is run in a separate process called the Kernel. The Kernel can be interrupted or restarted.

Practice: Execute WordCount in the Notebook (cluster)

+ Execute the WordCount spark application in the notebook

- Open the web navigator with the following url

<http://pirgi.udl.cat:8000>

- Login with your cluster account.
- Create a new pyspark notebook

- Create the Spark Context:

```
import pyspark  
sc = pyspark.SparkContext('local[*]')
```

- Copy the word count code, using the linux or hdfs file system for the input and output files.
- Execute (Ctrl+Enter or Menu Cell → Run Cells)

Standalone Spark Applications

- + The *spark-submit* script in Spark's bin directory is used to launch applications on a cluster.
 - It can use all of Spark's supported cluster managers (standalone, apache Mesos, Hadoop YARN) through a uniform interface.
- + Spark-submit syntax:

```
./bin/spark-submit --class <main-class> --master <master-url> \  
--deploy-mode <deploy-mode> --conf <key>=<value> \  
... # other options  
<application-jar> [application-arguments]
```
- + Bundling Your Application's Dependencies:
 - If your code depends on other projects, you will need to package them alongside your application in order to distribute the code to a Spark cluster (create an assembly jar).
 - For Python, you can use the `--py-files` argument of `spark-submit` to add `.py`, `.zip` or `.egg` files to be distributed with your application.

Example: Standalone Spark Applications

Run a Python application on a Spark standalone cluster

```
./bin/spark-submit --master spark://207.184.161.138:7077 \  
examples/src/main/python/pi.py 1000
```

Run on a YARN cluster

```
./bin/spark-submit --class org.apache.spark.examples.SparkPi \  
--master yarn --deploy-mode cluster \ # can be client for client mode \  
--executor-memory 20G --num-executors 50 \  
/path/to/examples.jar 1000
```

Run on our YARN cluster, with resource limitations.

```
export SPARK_HOME=/usr/hdp/2.3.4.0-3485/spark
```

```
export PYTHONPATH=$SPARK_HOME/python/:$PYTHONPATH
```

```
spark-submit --class org.apache.spark.examples.SparkPi --master yarn \  
--deploy-mode cluster --driver-memory 1g --executor-memory 512m \  
--executor-cores 1 $SPARK_HOME/lib/spark-examples*.jar 10
```

Practice: Execute WordCount in the standalone mode

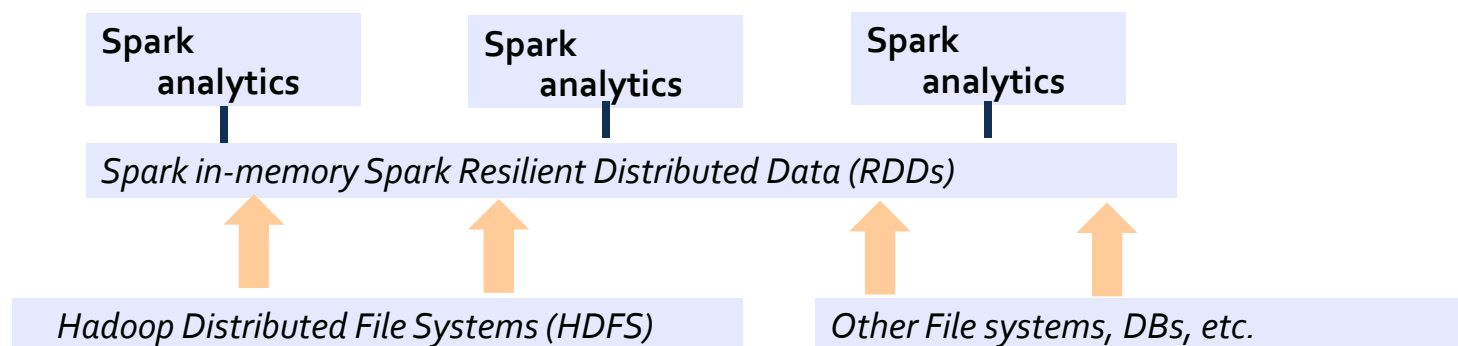
- + Execute the WordCount spark application in the standalone mode
 - Implement your spark+python program (wordcount example)
 - Copy the python file to the hadoop cluster
 - Submit the spark job:
 - `spark-submit --master yarn --deploy-mode cluster ./WordCount.py`
 - Check the results:
 - `hdfs dfs -cat <out_dir>/p*t`
 - See the output logs in http://pirgi.udl.cat:8088/proxy/application_1464002881097_0061/
- + Monitor your job
 - <http://pirgi.udl.cat:18081/>
- + On Error, consult the logs:
 - `yarn logs -applicationId application_1456747411295_0025`



Resilient Distributed Datasets (RDDs)


Resilient Distributed Datasets (RDDs)

- + Spark's basic unit of data
- + Immutable, fault tolerant collection of records that can be distributed and operated on in parallel across a cluster
- + Fault tolerance
 - If data in memory is lost it will be recreated from lineage
- + Caching, persistence (memory, spilling, disk) and check-pointing
- + Many database or file type can be supported



Resilient Distributed Datasets (RDDs)

- + An RDD is physically distributed across the cluster, but manipulated as one logical entity:
 - Each RDD is split into multiple **partitions**, which may be computed on different nodes of the cluster.
 - Spark will “distribute” any required processing to all partitions where the RDD exists and perform necessary redistributions and aggregations as well.
 - Partitioning can be based on a key in each record (using hash or range partitioning)
- + Example: Consider a distributed RDD “Names” made of names



Benefits of RDD Model

- + Consistency is easy due to immutability
- + Inexpensive fault tolerance (log lineage rather than replicating/checkpointing data)
- + Locality-aware scheduling of tasks on partitions
- + Despite being restricted, model seems applicable to a broad variety of applications

Creating RDDs

+ There are two ways to create RDDs:

- Parallelizing an existing collection in your driver program.

```
data = [1, 2, 3, 4, 5]  
distData = sc.parallelize(data)
```

- Referencing a dataset in an external storage system, such as a shared filesystem, HDFS, HBase, or any data source offering a Hadoop InputFormat.

```
distFile = sc.textFile("data.txt")
```

+ Once created, the distributed dataset can be operated on in parallel.

```
distData = distData.reduce(lambda a,b, a+b)
```

+ The number of partitions can be set manually by passing it as a second parameter to parallelize:

```
distData = sc.parallelize(data, 10)
```

- Spark will run one task for each partition of the cluster.
- Spark tries to set the number of partitions automatically based on your cluster.

External Datasets

- + Spark can create distributed datasets from any storage source supported by Hadoop:
 - Local filesystem, HDFS, Cassandra, HBase, Amazon S3, etc.
- + Using a path on the local filesystem, the file must also be accessible at the same path on worker nodes.
- + Support running on directories, compressed files, and wildcards
- + Apart from text files, Spark's supports several other data formats:
 - SparkContext.wholeTextFiles(path, minPart=None, use_unicode=True) lets read a directory containing multiple small text files, and returns each of them as (filename, content) pairs.
 - RDD.saveAsPickleFile(path, batchSize=10) and SparkContext.pickleFile(name, minPartitions=None) support saving an RDD as a SequenceFile of serialized pickled Python objects.
 - SequenceFile and Hadoop Input/Output Formats

RDD Operations

- + RDDs support two types of operations:
 - **Transformations:** which create a new dataset from an existing one.
 - *map*: is a transformation that passes each dataset element through a function and returns a new RDD representing the results
 - **Actions:** which return a value to the driver program after running a computation on the dataset.
 - *reduce*: is an action that aggregates all the elements of the RDD using some function and returns the final result to the driver program
- + Transformations in Spark are *lazy*, they do not compute their results right away.
 - The transformations are only computed when an action requires a result to be returned to the driver program.
- + Transformed RDD may be recomputed each time you run an action on it
 - A RDD may also *persist* in memory using the *persist* (or *cache*) method, to keep the elements around on the cluster to improve the access time.

Example: Transformations & Actions

```
text_file = sc.textFile("hdfs:///user/nando/data/books/quijote.txt")

# Calculate the length for each file line with map transformation
lineLengths = text_file.map(lambda s: len(s))

# Set lineLengths RDD to be keep in memory (persist)
print lineLengths.persist().is_cached

# Calculate the file length using the reduce action
totalLength = lineLengths.reduce(lambda a, b: a + b)
print "Total file length: "+format(totalLength)
```

RDD Operations

Transformations (define a new RDD)

map
filter
sample
union
groupByKey
reduceByKey
join
cache
...

Parallel operations/actions (return a result to driver)

reduce
collect
count
save
lookupKey
...



Programming with Spark

Passing Functions to Spark

+ Spark's API relies heavily on passing functions in the driver program to run on the cluster.

+ There are three recommended ways to do this:

- **Lambda expressions**, for simple functions that can be written as an expression.

```
field = self.field  
rdd.map(lambda s: field + s)
```

- **Local defs** inside the function calling into Spark, for longer code.

```
if __name__ == "__main__":  
    def myFunc(s):  
        words = s.split(" ")  
        return len(words)  
  
    sc = SparkContext(...)  
    sc.textFile("file.txt").map(myFunc)
```

- **Top-level functions** in a module.

Working with Key-Value Pairs

- + A few special operations are only available on RDDs of key-value pairs.
 - The most common ones are distributed “shuffle” operations, such as grouping or aggregating the elements by a key.
- + In Python, these operations work on RDDs containing built-in Python tuples such as (1, 2).
 - Simply create such tuples and then call your desired operation.
- + Example:

```
lines = sc.textFile("hdfs:///user/nando/data/books/quijote.txt")
pairs = lines.flatMap(lambda line: line.split(" ")).map(lambda w: (w, 1))
counts = pairs.reduceByKey(lambda a, b: a + b)
counts.sortByKey()
counts.collect()
```

Transformations (I)

- + ***map(func)***

Return a new distributed dataset formed by passing each element of the source through a function *func*.

- + ***filter(func)***

Return a new dataset formed by selecting those elements of the source on which *func* returns true.

- + ***flatMap(func)***

Similar to *map*, but each input item can be mapped to zero or more output items (so *func* should return a *Seq* rather than a single item).

- + ***mapPartitions(func)***

Similar to *map*, but runs separately on each partition (block) of the RDD, so *func* must be of type `Iterator<T> => Iterator<U>` when running on an RDD of type *T*.

Transformations (II)

- + **mapPartitionsWithIndex**(*func*)

Similar to mapPartitions, but also provides func with an integer value representing the index of the partition, so func must be of type (Int, Iterator<T>) => Iterator<U> when running on an RDD of type T.

- + **sample**(*withReplacement, fraction, seed*)

Sample a fraction of the data, with or without replacement, using a given random number generator seed.

- + **union**(*otherDataset*)

Return a new dataset that contains the union of the elements in the source dataset and the argument.

- + **intersection**(*otherDataset*)

Return a new RDD that contains the intersection of elements in the source dataset and the argument.

Transformations (III)

- + **distinct**(*[numPartitions=none]*)

Return a new dataset that contains the distinct elements of the source dataset.

- + **groupByKey**(*[numPartitions]*)

When called on a dataset of (K, V) pairs, returns a dataset of (K, Iterable<V>) pairs.

- + **reduceByKey**(*func, [numPartitions]*)

When called on a dataset of (K, V) pairs, returns a dataset of (K, V) pairs where the values for each key are aggregated using the given reduce function *func*, which must be of type $(V, V) \Rightarrow V$.

- + **aggregateByKey**(*zeroValue, seqOp, combOp, [numPartitions]*)

When called on a dataset of (K, V) pairs, returns a dataset of (K, U) pairs where the values for each key are aggregated using the given combine functions and a neutral "zero" value. Allows an aggregated value type that is different than the input value type, while avoiding unnecessary allocations.

Transformations (IV)

- + **sortByKey**(*[ascending]*, [NumPartitions])

When called on a dataset of (K, V) pairs where K implements Ordered, returns a dataset of (K, V) pairs sorted by keys in ascending or descending order, as specified in the boolean ascending argument.

- + **join**(*otherDataset*, [NumPartitions])

When called on datasets of type (K, V) and (K, W), returns a dataset of (K, (V, W)) pairs with all pairs of elements for each key.

- + **cogroup**(*otherDataset*, [NumPartitions])

When called on datasets of type (K, V) and (K, W), returns a dataset of (K, (Iterable<V>, Iterable<W>)) tuples. This operation is also called groupWith.

- + **cartesian**(*otherDataset*)

When called on datasets of types T and U, returns a dataset of (T, U) pairs (all pairs of elements).

Transformations (IV)

- + **pipe**(*command*, [*envVars*])

Pipe each partition of the RDD through a shell command, e.g. a Perl or bash script. RDD elements are written to the process's stdin and lines output to its stdout are returned as an RDD of strings.

- + **coalesce**(*numPartitions*)

Decrease the number of partitions in the RDD to *numPartitions*. Useful for running operations more efficiently after filtering down a large dataset.

- + **repartition**(*numPartitions*)

Reshuffle the data in the RDD randomly to create either more or fewer partitions and balance it across them. This always shuffles all data over the network.

- + **repartitionAndSortWithinPartitions**(*partitioner*)

Repartition the RDD according to the given partitioner and, within each resulting partition, sort records by their keys.

Actions (I)

- + **reduce(*func*)**

Aggregate the elements of the dataset using a function *func* (which takes two arguments and returns one). The function should be commutative and associative so that it can be computed correctly in parallel.

- + **collect()**

Return all the elements of the dataset as an array at the driver program. This is usually useful after a filter or other operation that returns a sufficiently small subset of the data.

- + **count()**

Return the number of elements in the dataset.

- + **first()**

Return the first element of the dataset (similar to `take(1)`).

- + **take(*n*)**

Return an array with the first *n* elements of the dataset.

Actions (II)

- + **takeSample**(*withReplacement*, *num*, [*seed*])

Return an array with a random sample of *num* elements of the dataset, with or without replacement, optionally pre-specifying a random number generator seed.

- + **takeOrdered**(*n*, [*ordering*])

Return the first *n* elements of the RDD using either their natural order or a custom comparator.

- + **saveAsTextFile**(*path*)

Write the elements of the dataset as a text file (or set of text files) in a given directory in the local filesystem, HDFS or any other Hadoop-supported file system.

- + **countByKey**()

Only available on RDDs of type (K, V). Returns a hashmap of (K, Int) pairs with the count of each key.

- + **foreach**(*func*)

Run a function *func* on each element of the dataset. This is usually done for side effects such as updating an [Accumulator](#) or interacting with external storage systems.



Questions???